

# GPU Applications for Modern Large Scale Asset Management

GTC 2014 – San José, California

Dr. Daniel Egloff

QuantAlea & IncubeAdvisory

March 27, 2014



## Portfolio Construction

## Portfolio Construction

### Constructing Asset Return Distributions

## Portfolio Construction

### Constructing Asset Return Distributions

#### I) Entropy Approach

Large Scale Convex Optimization with GPUs

## Portfolio Construction

### Constructing Asset Return Distributions

#### I) Entropy Approach

Large Scale Convex Optimization with GPUs

#### II) Bayesian Approach

Parallel Metropolis Hastings on GPUs

## Portfolio Construction

## Constructing Asset Return Distributions

### I) Entropy Approach

Large Scale Convex Optimization with GPUs

### II) Bayesian Approach

Parallel Metropolis Hastings on GPUs

## Conclusion

- ▶ Markowitz mean variance portfolio optimization
  - ▶ Optimally balance risk and performance

$$w_{\lambda}^* = \arg \max_{w, \mathbf{1}^T w = b} U_{\lambda}(w), \quad U_{\lambda}(w) = \mu^T w - \lambda w^T \Sigma w \quad (1)$$

- ▶ Risk budgeting
  - ▶ Construct portfolio so that risk contributions satisfy  $RC_i(w) = b_i$  for given risk budgets  $b_i$
  - ▶ For volatility risk measure we obtain

$$RC_i(w) = w_i \frac{(\Sigma w)_i}{\sqrt{w^T \Sigma w}} \quad (2)$$

- ▶ For both approaches we need means  $\mu$  and covariances  $\Sigma$  of asset returns

## Two steps

1. Prior asset return distribution from historical data or inverse optimization principle
2. Posterior asset return distribution should reflect expert views and beliefs
  - ▶ Views are distortions of a prior distribution
  - ▶ Views translate into constraints
  - ▶ Linear constraints on mean return pioneered by Black and Litterman

## We discuss two approaches

- ▶ Entropy based approach
- ▶ Bayesian methods



- ▶ Minimum discrimination principle (Kullback and Leibler 1951) infers a posterior distribution such that it
  - ▶ Satisfies a set of moments constraints
  - ▶ Is as hard as possible to differentiate from given prior distribution
- ▶ Minimum discrimination principle suitable methodology to incorporate constraints on a prior distribution which are
  - ▶ Linear
  - ▶ Quadratic
  - ▶ Convex

- ▶ Sample the probability distributions
  - ▶ Discrete sample space  $\mathcal{X} = \{x_1, \dots, x_m\}$
  - ▶ Probabilities elements of the standard simplex

$$p = (p_1, \dots, p_m)^\top \in \Delta^m \subset \mathbb{R}^m$$

- ▶ Implement views as distortions of prior probabilities  $p_i$
- ▶ Express views as feature functions  $f_j : \mathcal{X} \rightarrow \mathbb{R}, j = 1, \dots, n$
- ▶ Identify feature vector  $\mathbf{f} = (f_1, \dots, f_n)^\top : \mathcal{X} \rightarrow \mathbb{R}^n$  with matrix  $F = (f_{ji}) \in \mathbb{R}^{n \times m}$ , with  $f_{ji} = f_j(x_i)$



- ▶ Expectation of feature  $f : \mathcal{X} \rightarrow \mathbb{R}$  w.r.t measure  $q \in \Delta^m$  is

$$E_q[f] = \sum q_i f(x_i) = q^\top f$$

- ▶ Minimum discrimination principle defines probability  $p^*$  as

$$p^* = \arg \min_{q \in \Delta^m} q^\top (\log(q) - \log(p)) \quad (9)$$

subject to constraints

$$\mathbb{E}_q[f_j] \leq c_j, j = 1, \dots, n, \quad \mathbb{E}_q[\tilde{f}_j] = \tilde{c}_j, j = 1, \dots, \tilde{n}. \quad (10)$$

- ▶ Expectation views (10) lead to linear constraints

$$p^* = \arg \min_{\substack{q \in \mathbb{R}_+^m, \mathbf{1}^\top q = 1 \\ Fq \leq c, \tilde{F}q = \tilde{c}}} q^\top (\log(q) - \log(p)) \quad (11)$$

- ▶ Need large number of samples for accurate representation of distributions with  $m \simeq 10^6$
- ▶ With linear constraints: Lagrange duality to solve the optimal solution
  - ▶ Strong duality, i.e. no duality gap
  - ▶ Converting problem of dimension  $m$  to much smaller problems of dimension  $n + \tilde{n}$  to find optimal Lagrange multipliers

- ▶ How to extend to variance and covariance constraints?
- ▶ Lower variance constraint  $q^\top f^2 - (q^\top f)^2 \geq c$  convex but dual problem is ill-conditioned
- ▶ Upper variance constraint  $q^\top f^2 - (q^\top f)^2 \leq c$  not even convex
- ▶ Approximation by linearization  $q^\top f^2 \leq c + (p^\top f)^2$  to get a second moment constraint
- ▶ But .... linearized version often not accurate enough



- ▶ For convex constraints solve directly the primal problem
- ▶ Several accelerated first order methods (FOM) for convex optimization algorithms have been developed
- ▶ Modern algorithms adjust to the problem's geometry
- ▶ Examples
  - ▶ Mirror descent (MD) algorithm
  - ▶ Level bundle methods
  - ▶ Bundle mirror descent
  - ▶ Many different variations

- ▶ Gradient methods regularized with quadratic proximity term

$$y = \arg \min_{x \in \mathcal{X}} \{ \gamma \langle \xi, \nabla f(x) \rangle + c \|\xi - x\|^2 \}$$

- ▶ Replace  $\|\cdot\|^2$  with some distance like function  $D(\cdot, \cdot)$  that better exploits the geometry of  $\mathcal{X}$  and is simple to calculate
- ▶ Candidates: Bregman type distance based on kernel  
 $\omega : \mathcal{X} \rightarrow \mathbb{R}$

$$D(x, y) = D_\omega(x, y) \equiv \omega(x) - \omega(y) - \langle x - y, \omega'(y) \rangle$$

- ▶ Using  $\omega(u) = \frac{1}{2}\|u\|^2$  gives Euclidian norm distance

$$D_{\omega}(x, y) = \frac{1}{2}\|x - y\|^2$$

- ▶ The entropy  $\omega(x) = x^{\top} \log(x)$  on  $\mathcal{X} = \Delta^m$  leads to

$$D_{\omega}(x, y) = x^{\top} \log \left( \frac{x}{y} \right)$$

which is 1-strongly convex with respect to the  $L_1$  norm

$$D_{\omega}(x, y) \geq \frac{1}{2}\|x - y\|_1^2 \quad \forall x, y \in \Delta^m$$





- ▶ Convex problem  $\min_{x \in \mathcal{X}} f_0(x)$  subject to  $f_1(x) \leq 0$
- ▶  $\mathcal{X} \subset E$ , closed convex subset of f.d. vector space  $E$
- ▶  $\|\cdot\|$  norm on  $E$ ,  $\|\cdot\|_*$  dual norm on  $E^*$
- ▶ Duality pairing  $\langle \cdot, \cdot \rangle : E^* \times E \rightarrow \mathbb{R}$
- ▶ Bregman distance  $D_\omega(x, y)$  from kernel  $\omega : \mathcal{X} \rightarrow \mathbb{R}$
- ▶ Prox mapping  $\text{Prox}_x : E^* \rightarrow \mathcal{X}^\circ$

$$\text{Prox}_x(\xi) = \arg \min_{u \in \mathcal{X}} \{ \langle \xi, u \rangle + D_\omega(u, x) \}$$

- ▶ Mirror descent with constraints is a simple FOM
- ▶ Initial search point  $x_1 = x_\omega = \arg \min_{x \in \mathcal{X}} \omega(x)$
- ▶ For  $t = 1, \dots, N$  do
  - ▶ If  $f_1(x_t) \leq \text{tol}$  then  $i(t) = 0$  else  $i(t) = 1$
  - ▶ Select step size  $\gamma_t$
  - ▶ Update search point  $x_{t+1} = \text{Prox}_{x_t}(\gamma_t f'_{i(t)}(x_t))$
- ▶ Approximate solution after  $N$  steps

$$\hat{x}_N = \arg \min_{x \in \{x_t | i(t)=0\}} f_0(x)$$

- ▶ Possible choice  $\text{tol} = \gamma \|f'_1(x_t)\|_*$  and  $\gamma_t = \gamma \|f'_{i(t)}(x_t)\|_*^{-1}$

- ▶ Optimization problem (11) formulated on simplex  $\mathcal{X} = \Delta^m$
- ▶ Norm  $\|\cdot\|_1$  with dual norm  $\|\cdot\|_\infty$
- ▶ Prox mapping allows analytical expression

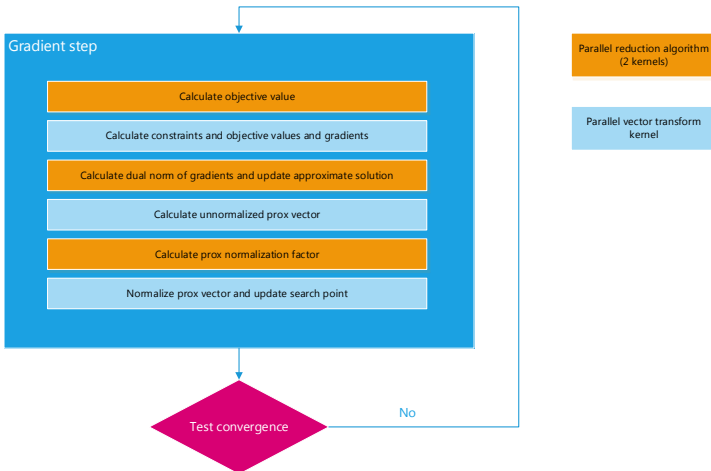
$$\text{Prox}_x(\xi) = \left(x^\top e^{-\xi}\right)^{-1} x \bullet e^{-\xi}$$

- ▶ Initial point  $x_\omega = n^{-1}\mathbf{1}_n$  the uniform distribution

- ▶ GPU implementation of constrained mirror descent algorithms for very large  $n$ 
  - ▶ Parallel coordinate-wise calculation of gradients for objective function and constraints
  - ▶ Vector reductions to calculate objective value, norms and scalar products
  - ▶ Parallel coordinate-wise calculation of prox vector, search points and execution of gradient step
- ▶ Gradient steps can be further parallelized by coordinate aggregation or randomization methods



# Schematic Gradient Step



- ▶ Usually kernel launch time ( $\approx 50\mu s$ ) not an issue
- ▶ For algorithms with many iterations, this becomes crucial
- ▶ Use kernel fusion technique to reduce number of kernel calls per gradient step
- ▶ Minimize number of arguments to pass to each kernel

- ▶ Bayesian methods treat the model parameters as random variables
- ▶ Views go into the prior distributions of the model parameters
- ▶ Historical data is used to update the prior distribution
- ▶ Not as flexible as entropy approach
- ▶ Does not rely on optimization hence less fragile
- ▶ Work horse is a parallel Metropolis Hastings Markov chain Monte Carlo sampler

- ▶ Parallelization of Metropolis Hastings through multiple parallel chains
- ▶ Further parallelization through pre-fetching
- ▶ Require efficient parallel random number generation and branch free methods to sample from proposal distribution



## GPU Metropolis Hastings sampler generation from generic target and proposal distribution expressions with F# and Alea.cuBase

```
let inline mhsamples (targetpdf : Expr<'T -> 'T>)
    (proppdf : Expr<'T -> 'T -> 'T -> 'T>)
    (proprnd : Expr<'T -> 'T -> 'T -> 'T>) = cuda {

    let! proppdf = proppdf |> Compiler.DefineInlineFunction
    let! proprnd = proprnd |> Compiler.DefineInlineFunction
    let! targetpdf = targetpdf |> Compiler.DefineInlineFunction

    let! kernel =
        <@ fun (steps:int) (scale:'T) (initial:deviceptr<'T>)
            (numbers:deviceptr<'T>) (result:deviceptr<'T>) ->
            let idx = blockIdx.x * blockDim.x + threadIdx.x
            let nchains = gridDim.x * blockDim.x

        . . .
```

Expressions for target distribution and proposal distribution are compiled directly into kernel

```
/// Dynamically generate expressions for MH at runtime
let targetpdf =
    <@ fun x ->
        exp(-x*x)*(2G + sin(5G*x) + sin(2G*x)) @>
let proppdf =
    <@ fun sigma x y ->
        exp(-(y - x)*(y - x)/(2G*sigma*sigma)) / (sigma*__sqrt2pi()) @>
let proprng =
    <@ fun sigma x u -> x + sigma*u @>

/// Compiler integrated into F#, callable through API
let mh = mhsamples targetpdf proppdf proprng |> Util.load Worker.Default
```




Figure : Sample histogram for target density  
 $e^{-x^2}(2 + \sin(5x) + \sin(2x))$




Figure : Kernel density smoothed sample pdf for target density  
 $e^{-x^2}(2 + \sin(5x) + \sin(2x))$

- ▶ We have shown two applications of GPUs for quantitative asset management
- ▶ More applications in strategy calibration, back-testing and MIP
- ▶ Large scale convex optimization occurs in other fields such as
  - ▶ Machine learning such as e.g. deep believe network training in a big data context
  - ▶ Network analysis
  - ▶ Truss topology design
  - ▶ Compressed sensing
- ▶ Markov chain Monte Carlo methods require flexible ways to specify various parts such as target pdf, the proposal distribution and a dynamic way to derive tuning parameters, for which our Alea.cuBase F# GPU compiler is ideal